

- partn_nyc index is made on the part_number column in nyc because in my queries I retrieve a ID's from the table more than once and the table is quite large.
- partn_sfo is virtually the same index as partn_nyc except that the index is made using the sfo table and not the nyc. Both of these indexes helped with my Third query where I had to retrieve subqueries of all the part numbers for each table. Indexes partn_nyc and partn_sfo reduced the run time for query 2 from 23.956 ms to 20.588 ms
- supld, sfo_sup, nyc_sup were created because of how much interaction there is between the supplier_id column in suppliers and the supplier column in both part_nyc and part_sfo. The index would reduce how many times I would have to traverse the whole data set to get the subqueries where they interact. In query 3 the indexes listed above reduced the runtime from 43.655 ms to 40.241 ms.
- I did not index onhand at all since this attribute is changed in query 5 and 6 and this would cause an expensive update.